

Register Allocation

Lecture 16

Prof. Aiken CS 143 Lecture 16

1





Lecture Outline

- Memory Hierarchy Management
- Register Allocation
 - Register interference graph
 - Graph coloring heuristics
 - Spilling
- Cache Management

Prof. Aiken CS 143 Lecture 16

2

The Memory Hierarchy

	Registers	1 cycle	256-8000 bytes
	Cache	3 cycles	256k-1M
	Main memory	20-100 cycles	32M-1G
	Disk	0.5-5M cycles	4G-1T

Prof. Aiken CS 143 Lecture 16

3

Managing the Memory Hierarchy

- Programs are written as if there are only two kinds of memory: main memory and disk
 - Programmer is responsible for moving data from disk to memory (e.g., file I/O)
 - Hardware is responsible for moving data between memory and caches
 - Compiler is responsible for moving data between memory and registers

Prof. Aiken CS 143 Lecture 16

4

Current Trends

- Power usage limits
 - Size and speed of registers/caches
 - Speed of processors
- But
 - The cost of a cache miss is very high
 - Typically requires 2 caches to bridge fast processor with large main memory
- It is very important to:
 - Manage registers properly
 - Manage caches properly
- Compilers are good at managing registers

Prof. Aiken CS 143 Lecture 16

5

The Register Allocation Problem

- Intermediate code uses unlimited temporaries
 - Simplifies code generation and optimization
 - Complicates final translation to assembly
- Typical intermediate code uses too many temporaries

Prof. Aiken CS 143 Lecture 16

6

The Register Allocation Problem (Cont.)

- The problem:
Rewrite the intermediate code to use no more temporaries than there are machine registers
- Method:
 - Assign multiple temporaries to each register
 - But without changing the program behavior

Prof. Aiken CS 143 Lecture 16

7

History

- Register allocation is as old as compilers
 - Register allocation was used in the original FORTRAN compiler in the '50s
 - Very crude algorithms
- A breakthrough came in 1980
 - Register allocation scheme based on graph coloring
 - Relatively simple, global and works well in practice

Prof. Aiken CS 143 Lecture 16

8

An Example

- Consider the program
 - $a := c + d$
 - $e := a + b$
 - $f := e - 1$
- Can allocate a , e , and f all to one register (r_1):
 - $r_1 := r_2 + r_3$
 - $r_1 := r_1 + r_4$
 - $r_1 := r_1 - 1$
- Assume a and e dead after use
 - Temporary a can be "reused" after $e := a + b$
 - So can temporary e
- A dead temporary is not needed
 - A dead temporary can be reused

Prof. Aiken CS 143 Lecture 16

9

The Idea

Temporaries t_1 and t_2 can share the same register if at any point in the program at most one of t_1 or t_2 is live.

Or

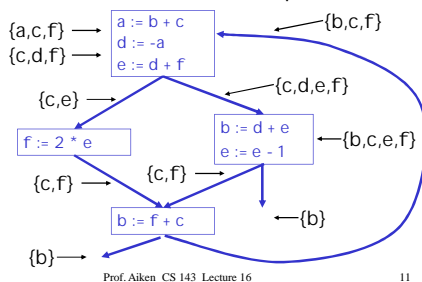
If t_1 and t_2 are live at the same time, they cannot share a register

Prof. Aiken CS 143 Lecture 16

10

Algorithm: Part I

- Compute live variables for each point:



Prof. Aiken CS 143 Lecture 16

11

The Register Interference Graph


- Construct an undirected graph
 - A node for each temporary
 - An edge between t_1 and t_2 if they are live simultaneously at some point in the program
- This is the *register interference graph* (RIG)
 - Two temporaries can be allocated to the same register if there is no edge connecting them

Prof. Aiken CS 143 Lecture 16

12

Example

- For our example:



- E.g., b and c cannot be in the same register
- E.g., b and d could be in the same register

Prof. Aiken CS 143 Lecture 16

13

Notes on Register Interference Graphs

- Extracts exactly the information needed to characterize legal register assignments
- Gives a global (i.e., over the entire flow graph) picture of the register requirements
- After RIG construction the register allocation algorithm is architecture independent

Prof. Aiken CS 143 Lecture 16

14

Definitions

- A coloring of a graph is an assignment of colors to nodes, such that nodes connected by an edge have different colors
- A graph is k-colorable if it has a coloring with k colors

Prof. Aiken CS 143 Lecture 16

15

Register Allocation Through Graph Coloring


- In our problem, colors = registers
 - We need to assign colors (registers) to graph nodes (temporaries)
- Let k = number of machine registers
- If the RIG is k -colorable then there is a register assignment that uses no more than k registers

Prof. Aiken CS 143 Lecture 16

16

Graph Coloring Example

- Consider the example RIG




- There is no coloring with less than 4 colors
- There are 4-colorings of this graph

Prof. Aiken CS 143 Lecture 16

17

Example Review




Prof. Aiken CS 143 Lecture 16

18

Example After Register Allocation

- Under this coloring the code becomes:



Prof. Aiken CS 143 Lecture 16

19

Computing Graph Colorings

- How do we compute graph colorings?
- It isn't easy:
 - This problem is very hard (NP-hard). No efficient algorithms are known.
 - Solution: use heuristics
 - A coloring might not exist for a given number of registers
 - Solution: later

Prof. Aiken CS 143 Lecture 16

20

Graph Coloring Heuristic

- Observation:
 - Pick a node t with fewer than k neighbors in RIG
 - Eliminate t and its edges from RIG
 - If resulting graph is k -colorable, then so is the original graph
- Why?
 - Let c_1, \dots, c_n be the colors assigned to the neighbors of t in the reduced graph
 - Since $n < k$ we can pick some color for t that is different from those of its neighbors

Prof. Aiken CS 143 Lecture 16

21

Graph Coloring Heuristic


- The following works well in practice:
 - Pick a node t with fewer than k neighbors
 - Put t on a stack and remove it from the RIG
 - Repeat until the graph has one node
- Assign colors to nodes on the stack
 - Start with the last node added
 - At each step pick a color different from those assigned to already colored neighbors

Prof. Aiken CS 143 Lecture 16

22

Graph Coloring Example (1)

- Start with the RIG and with $k = 4$:




- Remove a

Prof. Aiken CS 143 Lecture 16

23

Graph Coloring Example (2)




- Remove d

Prof. Aiken CS 143 Lecture 16

24

Graph Coloring Example (3)

- Note: all nodes now have fewer than 4 neighbors




Stack: {d, a}

- Remove c

Prof. Aiken CS 143 Lecture 16

25

Graph Coloring Example (4)



Stack: {c, d, a}

- Remove b

Prof. Aiken CS 143 Lecture 16

26

Graph Coloring Example (5)



Stack: {b, c, d, a}

- Remove e

Prof. Aiken CS 143 Lecture 16

27

Graph Coloring Example (6)



Stack: {e, b, c, d, a}

- Remove f

Prof. Aiken CS 143 Lecture 16

28

Graph Coloring Example (7)

- Now start assigning colors to nodes, starting with the top of the stack

Stack: {f, e, b, c, d, a}

Prof. Aiken CS 143 Lecture 16

29

Graph Coloring Example (8)



Stack: {e, b, c, d, a}

Prof. Aiken CS 143 Lecture 16

30

Graph Coloring Example (9)




Stack: {b, c, d, a}

- e must be in a different register from f

Prof. Aiken CS 143 Lecture 16

31

Graph Coloring Example (10)




Stack: {c, d, a}

Prof. Aiken CS 143 Lecture 16

32

Graph Coloring Example (11)




Stack: {d, a}

Prof. Aiken CS 143 Lecture 16

33

Graph Coloring Example (12)




Stack: {a}

Prof. Aiken CS 143 Lecture 16

34

- d can be in the same register as b

Graph Coloring Example (13)




Prof. Aiken CS 143 Lecture 16

35

What if the Heuristic Fails?

- What if all nodes have k or more neighbors ?
- Example: Try to find a 3-coloring of the RIG:




Prof. Aiken CS 143 Lecture 16

36

What if the Heuristic Fails?

- Remove a and get stuck (as shown below)
- Pick a node as a candidate for *spilling*
 - A spilled temporary "lives" in memory
 - Assume that f is picked as a candidate




Prof. Aiken CS 143 Lecture 16

37

What if the Heuristic Fails?

- Remove f and continue the simplification
 - Simplification now succeeds: b, d, e, c




Prof. Aiken CS 143 Lecture 16

38

What if the Heuristic Fails?

- Eventually we must assign a color to f
- We hope that among the 4 neighbors of f we use less than 3 colors \Rightarrow optimistic coloring



Prof. Aiken CS 143 Lecture 16

39

Spilling


- If optimistic coloring fails, we spill f
 - Allocate a memory location for f
 - Typically in the current stack frame
 - Call this address fa
- Before each operation that reads f , insert $f := \text{load } fa$
- After each operation that writes f , insert $\text{store } f, fa$

Prof. Aiken CS 143 Lecture 16

40

Spilling Example

- This is the new code after spilling f

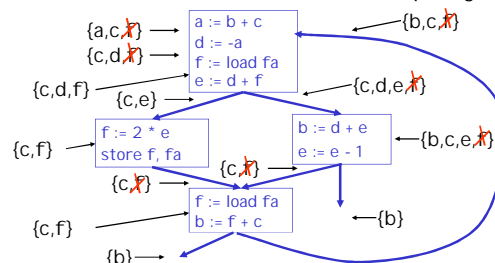


Prof. Aiken CS 143 Lecture 16

41

Recomputing Liveness Information

- The new liveness information after spilling:



Prof. Aiken CS 143 Lecture 16

42

Recomputing Liveness Information


- New liveness information is almost as before
- f is live only
 - Between a $f := \text{load } fa$ and the next instruction
 - Between a $\text{store } f, fa$ and the preceding instr.
- Spilling reduces the live range of f
 - And thus reduces its interferences
 - Which results in fewer RIG neighbors for f

Prof. Aiken CS 143 Lecture 16

43

Recompute RIG After Spilling

- Some edges of the spilled node are removed
- In our case f still interferes only with c and d
- And the resulting RIG is 3-colorable



Prof. Aiken CS 143 Lecture 16

44

Spilling Notes

- Additional spills might be required before a coloring is found
- The tricky part is deciding what to spill
 - But any choice is correct
- Possible heuristics:
 - Spill temporaries with most conflicts
 - Spill temporaries with few definitions and uses
 - Avoid spilling in inner loops

Prof. Aiken CS 143 Lecture 16

45

Caches

- Compilers are very good at managing registers
 - Much better than a programmer could be
- Compilers are not good at managing caches
 - This problem is still left to programmers
 - It is still an open question how much a compiler can do to improve cache performance
- Compilers can, and a few do, perform some cache optimizations

Prof. Aiken CS 143 Lecture 16

46

Cache Optimization

- Consider the loop

```
for(j := 1; j < 10; j++)
  for(i=1; i<1000; i++)
    a[i] *= b[i]
```
- This program has terrible cache performance
 - Why?

Prof. Aiken CS 143 Lecture 16

47

Cache Optimization (Cont.)

- Consider the program:

```
for(i=1; i<1000; i++)
  for(j := 1; j < 10; j++)
    a[i] *= b[j]
```

 - Computes the same thing
 - But with much better cache behavior
 - Might actually be more than 10x faster
- A compiler can perform this optimization
 - called loop interchange

Prof. Aiken CS 143 Lecture 16

48

Conclusions

- Register allocation is a “must have” in compilers:
 - Because intermediate code uses too many temporaries
 - Because it makes a big difference in performance
- Register allocation is more complicated for CISC machines