

Language Security

Lecture 19

Prof. Aiken CS 143 Lecture 19

1

Lecture Outline

- Beyond compilers
 - Looking at other issues in programming language design and tools
- C
 - Arrays
 - Exploiting buffer overruns
 - Detecting buffer overruns

Prof. Aiken CS 143 Lecture 19

2

Platitudes

- Language design has influence on
 - Safety
 - Efficiency
 - Security

Prof. Aiken CS 143 Lecture 19

3

C Design Principles

- Small language
- Maximum efficiency
- Safety less important
- Designed for the world in 1972
 - Weak machines
 - Trusted networks


Prof. Aiken CS 143 Lecture 19

4

Arrays in C

```
char buffer[100];
```

Declares and allocates an array of 100 chars



Prof. Aiken CS 143 Lecture 19

5

C Array Operations

```
char buf1[100], buf2[100];
```

Write:

```
buf1[0] = 'a';
```

Read:

```
return buf2[0];
```

Prof. Aiken CS 143 Lecture 19

6

What's Wrong with this Picture?

```
int i = 0;
for(i = 0; buf1[i] != '\0'; i++)
    { buf2[i] = buf1[i]; }
buf2[i] = '\0';
```

Prof. Aiken CS 143 Lecture 19

7

Indexing Out of Bounds

The following are all legal C and may generate no run-time errors

```
char buffer[100];

buffer[-1] = 'a';
buffer[100] = 'a';
buffer[100000] = 'a';
```

Prof. Aiken CS 143 Lecture 19

8

Why?

- Why does C allow out of bounds array references?
 - Proving at compile-time that all array references are in bounds is very difficult (impossible in C)
 - Checking at run-time that all array references are in bounds is expensive

Prof. Aiken CS 143 Lecture 19

9

Code Generation for Arrays

```
buf1[i] = 1; /* buf1 has type int[] */

r1 = load &buf1;
r2 = load i;
r3 = r2 * 4;
r4 = r1 + r3
store r4, 1
```

Prof. Aiken CS 143 Lecture 19

10

Discussion

- 5 instructions worst case
- Often `&buf1` and `i` already in registers
 - Saves 2 instructions
- Many machines have indirect loads/stores
 - `store r1[r3], 1`
 - Saves 1 instruction
- Best case 2 instructions
 - Offset calculation and memory operation

Prof. Aiken CS 143 Lecture 19

11

Code Generation for Arrays with Bounds Checks

```
buf1[i] = 1; /* buf1 has type int[] */

r1 = load &buf1;
r2 = load i;
r3 = r2 * 4;
if r3 < 0 then error;
r5 = load limit of buf1;
if r3 >= r5 then error;
r4 = r1 + r3
store r4, 1
```

Prof. Aiken CS 143 Lecture 19

12

Discussion

- Lower bounds check can often be removed
 - Easy to prove statically that index is positive
- Upper bounds check hard to remove
 - Leaves a conditional in instruction stream
- In C, array limits not stored with array
 - Knowing the array limit for a given reference is non-trivial

Prof. Aiken CS 143 Lecture 19

13

C vs. Java

- C array reference typical case
 - Offset calculation
 - Memory operation (load or store)
- Java array reference typical case
 - Offset calculation
 - Memory operation (load or store)
 - Array bounds check
 - Type compatibility check (for stores)

Prof. Aiken CS 143 Lecture 19

14

Buffer Overruns

- A buffer overrun writes past the end of an array
- *Buffer* usually refers to a C array of char
 - But can be any array
- So who's afraid of a buffer overrun?
 - Can damage data structures
 - Cause a core dump
 - What else?


Prof. Aiken CS 143 Lecture 19

15

Stack Smashing

Buffer overruns can alter the control flow of your program!

`char buffer[100]; /* stack allocated array */`



Prof. Aiken CS 143 Lecture 19

16

An Overrun Vulnerability



```
void foo(char buf1[]) {  
    char buf2[100];  
    int i = 0;  
    for(i = 0; buf1[i] != '\0'; i++)  
        { buf2[i] = buf1[i]; }  
    buf2[i] = '\0';  
}
```

Prof. Aiken CS 143 Lecture 19

17

An Interesting Idea

```
char buf[104] = { ' ', ..., ' ', magic 4 chars }  
foo(buf); (**)
```



Prof. Aiken CS 143 Lecture 19

18

Discussion

- So we can make `foo` jump wherever we like.
- How is this possible?
- Unanticipated interaction of two features:
 - Unchecked array operations
 - Stack-allocated arrays
 - Knowledge of frame layout allows prediction of where array and return address are stored
 - Note the “magic cast” from char’s to an address

Prof. Aiken CS 143 Lecture 19

19

The Rest of the Story


- We can make `foo` jump anywhere.
- But where is a useful place to jump?
- Idea: Put our own code in the buffer and jump there!

Prof. Aiken CS 143 Lecture 19

20

The Plan

```
char buf[104] = { 104 magic chars }  
foo(buf);
```



Prof. Aiken CS 143 Lecture 19

21

Details

- “exec /bin/sh”
 - Easy to write in assembly code
 - Make all jumps relative
- Be careful not to have null's in the code (why?)

Prof. Aiken CS 143 Lecture 19

22

More Details

- Overwrite return address with start of buffer
 - Harder
 - Need to guess where buffer in called routine starts (trial & error)
 - Pad front of buffer with NOPS
 - Guess need not be exact; just land somewhere in NOPS

Prof. Aiken CS 143 Lecture 19

23

And More Details

- Overwrite return address
 - Don't need to know exactly where return address is
 - Just pad end of buffer with multiple copies of new return address X

```
char buf[104] =  
    "NOPS ... /bin/exec sh XXXXXXXXXXXX"  
foo(buf);
```

Prof. Aiken CS 143 Lecture 19

24

The State of C Programming

- Buffer overruns are common
 - Programmers must do their own bounds checking
 - Easy to forget or be off-by-one or more
 - Program still appears to work correctly
- In C wrt to buffer overruns
 - Easy to do the wrong thing
 - Hard to do the right thing

Prof. Aiken CS 143 Lecture 19

25

The State of Hacking

- Buffer overruns are the attack of choice
 - 40-50% of new vulnerabilities are buffer overrun exploits
- Highly automated toolkits available to exploit known buffer overruns
 - Search for "buffer overruns" yields > 25,000 hits

Prof. Aiken CS 143 Lecture 19

26

The Sad Reality

- Even well-known buffer overruns are still widely exploited
 - Hard to get people to upgrade millions of vulnerable machines
- We assume that there are many more unknown buffer overrun vulnerabilities
 - At least unknown to the good guys

Prof. Aiken CS 143 Lecture 19

27

Static Analysis to Detect Buffer Overruns

- Detecting buffer overruns *before* distributing code would be better
- Idea: Build a tool similar to a type checker to detect buffer overruns
- Joint work with David Wagner & Jeff Foster

Prof. Aiken CS 143 Lecture 19

28

Focus on Strings

- Most important buffer overrun exploits are through string buffers
 - Reading an untrusted string from the network, keyboard, etc.
- Focus the tool only on arrays of characters

Prof. Aiken CS 143 Lecture 19

29

Idea 1: Strings as an Abstract Data Type

- A problem: Pointer operations & array dereferences are very difficult to analyze statically
 - Where does **a* point?
 - What does *buf[j]* refer to?
- Idea: Model effect of string library functions directly
 - Hard code effect of *strcpy*, *strcat*, etc.

Prof. Aiken CS 143 Lecture 19

30

Idea 2: The Abstraction

- Model buffers as pairs of integer ranges
 - *Size* allocated size of the buffer in bytes
 - *Length* number of bytes actually in use
- Use integer ranges $[x,y] = \{x, x+1, \dots, y-1, y\}$
 - Size & length cannot be computed exactly

Prof. Aiken CS 143 Lecture 19

31

The Strategy

- For each program expression, write constraints capturing the *alloc* and *len* of its string subexpressions
- Solve the constraints for the entire program
- Check for each string variable *s*
 $\text{len}(s) \leq \text{alloc}(s)$

Prof. Aiken CS 143 Lecture 19

32

The Constraints

<code>char s[n];</code>	$n \subseteq \text{alloc}(s)$
<code>strcpy(dst,src)</code>	$\text{len}(\text{src}) \subseteq \text{len}(\text{dst})$
<code>p = strdup(s)</code>	$\text{len}(s) \subseteq \text{len}(p) \ \& \ \text{alloc}(s) \subseteq \text{alloc}(p)$
<code>p[n] = '\0'</code>	$\min(\text{len}(p), n+1) \subseteq \text{len}(p)$

Prof. Aiken CS 143 Lecture 19

33

Constraint Solving

- Solving the constraints is akin to solving dataflow equations (e.g., constant propagation)
- Build a graph
 - Nodes are $\text{len}(s)$, $\text{alloc}(s)$
 - Edges are constraints $\text{len}(s) \subseteq \text{len}(t)$
- Propagate information forward through the graph
 - Special handling of loops in the graph

Prof. Aiken CS 143 Lecture 19

34

Results

- Found new buffer overruns in sendmail
- Found new exploitable overruns in Linux nettools package
- Both widely used, previously hand-audited packages

Prof. Aiken CS 143 Lecture 19

35

Limitations

- Tool produces many false positives
 - 1 out of 10 warnings is a real bug
- Tool has false negatives
 - Unsound---may miss some overruns
- Newer tools greatly improve on these results
 - E.g., METAL, Microsoft's SAL

Prof. Aiken CS 143 Lecture 19

36

Summary

- Programming language knowledge useful beyond compilers
- Useful for programmers
 - Understand what you are doing!
- Useful for tools other than compilers
 - Big research direction

The Last Slide

- Have a great New Year!